Minimal Audit Protocol

For LLM-sentiment return-predictability studies (one page, checklist form)

Companion to *Dissecting LLM News Sentiment Scores: De-Identification, Batch Composition, Scoring Noise, and Price Echoes in Return Predictability* (working paper, Appendix).

Replication package: https://github.com/43aquaris/dissecting-llm-news-sentiment-scores

---

The protocol below is the paper's working answer to "what should a small-sample LLM-sentiment study run before reporting a predictability statistic?" It is deliberately minimal: every component costs at most one additional scoring pass over the study's own corpus, and the whole protocol is executable in an afternoon at headline-scale prices. Each item names the failure mode it guards against, as documented somewhere in this paper.

1. **Anonymize once, and verify the anonymization.** Replace firm names and tickers with a neutral placeholder, then run a forced residual-identifier check over every transformed text. Guards against identity priors, and against silent no-op anonymizations (the ticker-V failure of Appendix **??**: ten headlines whose "anonymized" text still named the firm).

2. **Field the named-versus-anonymized contrast interleaved in a single sitting**, alternating condition order at the batch level. Guards against cross-session level drift masquerading as a treatment effect (+5.79 points over twelve hours here; Section **??**).

3. **Freeze the batch composition** across every comparison —conditions, sessions, or model versions —and re-score once under an alternative composition to measure the channel on your own corpus. Guards against batch-composition level shifts (+3.6 points here; harmless to slopes and ranks, disqualifying for thresholds and time series; Appendix **??**).

4. **Score twice and report the noise floor:** an ICC or within-headline repeat SD alongside any single-call statistic, with inference clustered at the call level. Guards against response-level mood and overstated precision (ICC 0.92 here; 13% of residual variance is shared within a response; Section **??**).

5. **Flag realized-price language and audit the flag** against a written rulebook, on every flagged headline if the count is small, on a sample otherwise. Guards against retrospective-coverage contamination (the echo channel of Section **??**: the predictive loading directionally sits in the 15% of headlines that quote realized moves).

6. **Test the sample for selection:** a two-line firm-and-day fixed-effects regression of realized activity on the coverage indicator. Guards against event-sampling bounds on external validity (coverage selected on volatility at $t = 2.76$ here; Section **??**).

7. **Declare the test family before the augmented inference, and compare instruments by difference tests, never by separate significances.** Guards against multiplicity and the genre's most common inference error (Sections **??** and **??**).

Items 1–4 protect the measurement; items 5–6 protect the interpretation of the return regression; item 7 protects the inference. A study that cannot run the full protocol should at minimum report which items it skipped, so that readers can price the corresponding channels themselves.

Checklist form of Appendix "A Minimal Viable Audit Protocol"; the paper's appendix is the authoritative text and the two are generated from the same source. Free to reuse with citation.